

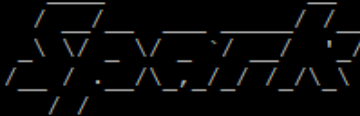
# Spark Hands-On

Dr. Villanes

Once you see this screen.. Type **pyspark**

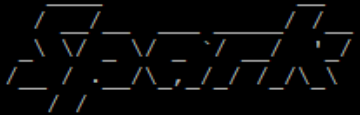
```
hadoop@ip-172-31-66-228:~  
  
  _I  ( _I_ )  
  _I  ( _I_ /  
  _I \ _I_ _I  
Amazon Linux 2 AMI  
  
https://aws.amazon.com/amazon-linux-2/  
16 package(s) needed for security, out of 51 available  
Run "sudo yum update" to apply all updates.  
  
EEEEEEEEEEEEEEEEEEEE MMMMMMMM MMMMMMMM RRRRRRRRRRRRRRRR  
E::::::::::::::::::E M::::::::M M::::::::M R::::::::::::R  
EE::::::::EEEEEEEEEE E M::::::::M M::::::::M R::::RRRRRR::::R  
E::::E EEEEE M::::::::M M::::::::M RR::::R R::::R  
E::::E M::::::::M M::M::::M M::M::::M R::R R::::R  
E::::EEEEEEEEEE M::::M M::M M::M M::::M R::RRRRRR::::R  
E::::::::::::::::::E M::::M M::M::::M M::::M R::::::::::::RR  
E::::EEEEEEEEEE M::::M M::::M M::::M R::RRRRRR::::R  
E::::E M::::M M::M M::::M R::R R::::R  
E::::E EEEEE M::::M MMM M::::M R::R R::::R  
EE::::::::EEEEEEEEEE E M::::M M::::M R::R R::::R  
E::::::::::::::::::E M::::M M::::M RR::::R R::::R  
EEEEEEEEEEEEEEEEEEEE MMMMMMMM MMMMMMMM RRRRRRR RRRRRR  
  
[hadoop@ip-172-31-66-228 ~]$ pyspark
```

You should now see something like this...

```
hadoop@ip-172-31-43-113:~  
Setting default log level to "WARN".  
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).  
24/02/08 16:16:50 WARN HiveConf: HiveConf of name hive.server2.thrift.url does not exist  
24/02/08 16:16:54 WARN Client: Neither spark.yarn.jars nor spark.yarn.archive is set, falling back to uploading libraries under SPARK_HOME.  
24/02/08 16:17:41 WARN YarnSchedulerBackend$YarnSchedulerEndpoint: Attempted to request executors before the AM has registered!  
Welcome to  
 version 3.5.0-amzn-0  
Using Python version 3.9.16 (main, Sep 8 2023 00:00:00)  
Spark context Web UI available at http://ip-172-31-43-113.ec2.internal:4040  
Spark context available as 'sc' (master = yarn, app id = application_1707408358037_0001).  
SparkSession available as 'spark'.  
>>> pyspark  
<module 'pyspark' from '/usr/lib/spark/python/pyspark/__init__.py'>  
>>> 
```

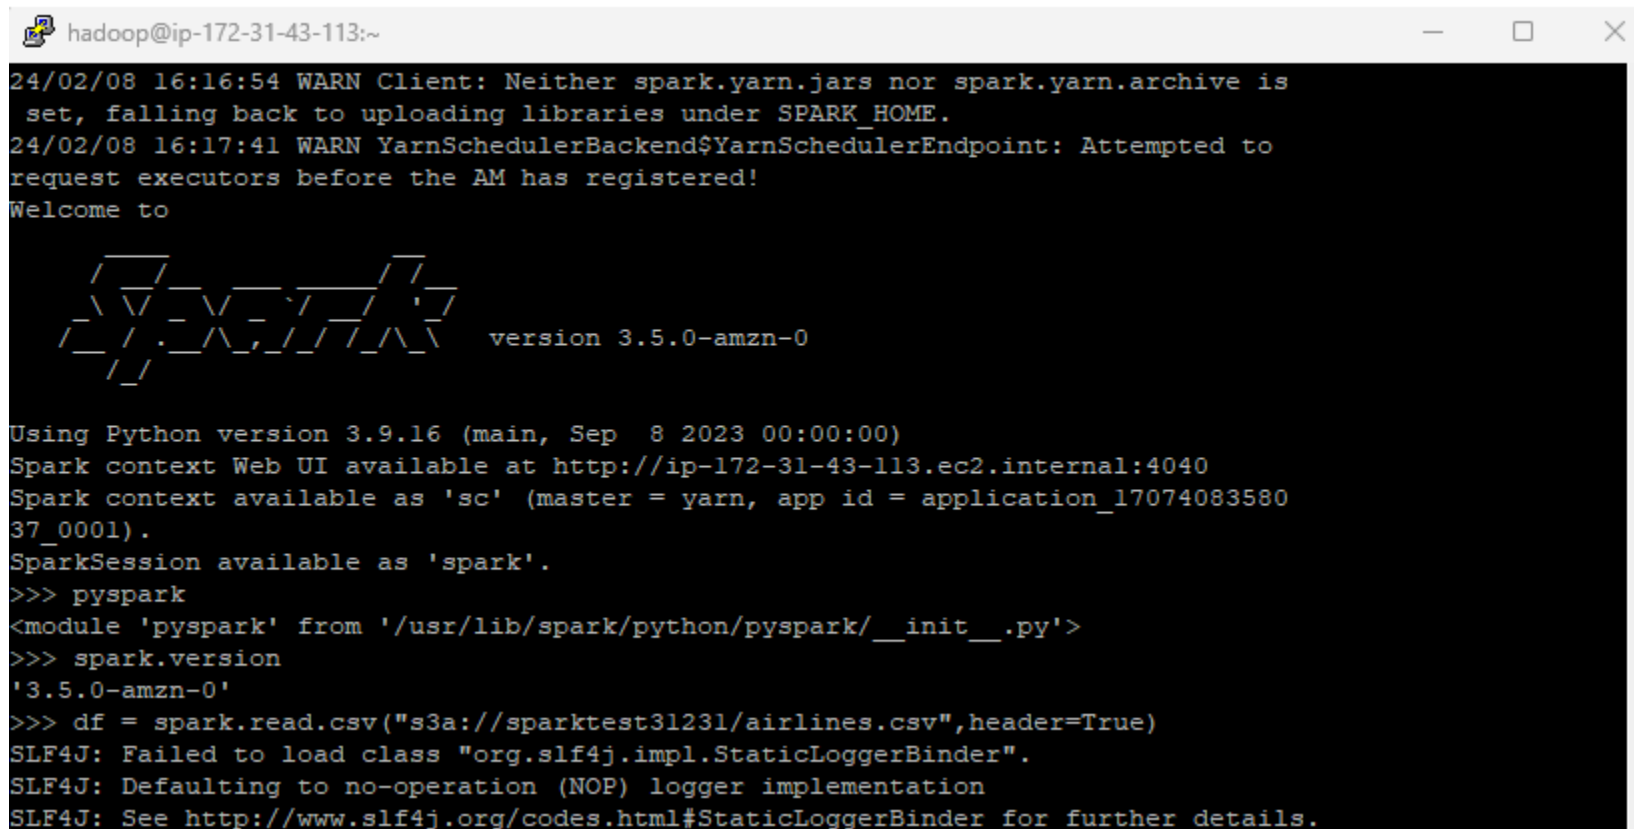
# Spark Version

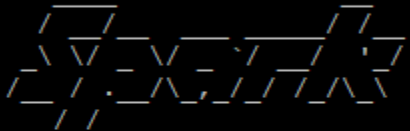
>>> spark.version

```
hadoop@ip-172-31-43-113:~  
1(newLevel) .  
24/02/08 16:16:50 WARN HiveConf: HiveConf of name hive.server2.thrift.url does n  
ot exist  
24/02/08 16:16:54 WARN Client: Neither spark.yarn.jars nor spark.yarn.archive is  
set, falling back to uploading libraries under SPARK_HOME.  
24/02/08 16:17:41 WARN YarnSchedulerBackend$YarnSchedulerEndpoint: Attempted to  
request executors before the AM has registered!  
Welcome to  
 version 3.5.0-amzn-0  
Using Python version 3.9.16 (main, Sep  8 2023 00:00:00)  
Spark context Web UI available at http://ip-172-31-43-113.ec2.internal:4040  
Spark context available as 'sc' (master = yarn, app id = application_17074083580  
37_0001).  
SparkSession available as 'spark'.  
>>> pyspark  
<module 'pyspark' from '/usr/lib/spark/python/pyspark/__init__.py'>  
>>> spark.version  
'3.5.0-amzn-0'  
>>> 
```

# Creating Spark Dataframe

```
>>> df = spark.read.csv("s3a://bucket_name/airlines.csv",header=True)
```

A terminal window titled 'hadoop@ip-172-31-43-113:~' showing the Spark shell startup process. It includes warning messages about missing jars and the Yarn scheduler endpoint, followed by a 'Welcome to' message and the Spark logo. The terminal then shows the Spark version (3.5.0-amzn-0), the Python version (3.9.16), and the Spark context Web UI URL. It also shows the Spark context available as 'sc' and the SparkSession available as 'spark'. The user enters 'pyspark' and 'spark.version', and then executes the code to read a CSV file from S3. The output shows the code execution and a warning message about the SLF4J logger implementation.

```
hadoop@ip-172-31-43-113:~  
24/02/08 16:16:54 WARN Client: Neither spark.yarn.jars nor spark.yarn.archive is  
set, falling back to uploading libraries under SPARK_HOME.  
24/02/08 16:17:41 WARN YarnSchedulerBackend$YarnSchedulerEndpoint: Attempted to  
request executors before the AM has registered!  
Welcome to  
 version 3.5.0-amzn-0  
Using Python version 3.9.16 (main, Sep 8 2023 00:00:00)  
Spark context Web UI available at http://ip-172-31-43-113.ec2.internal:4040  
Spark context available as 'sc' (master = yarn, app id = application_17074083580  
37_0001).  
SparkSession available as 'spark'.  
>>> pyspark  
<module 'pyspark' from '/usr/lib/spark/python/pyspark/__init__.py'>  
>>> spark.version  
'3.5.0-amzn-0'  
>>> df = spark.read.csv("s3a://sparktest31231/airlines.csv",header=True)  
SLF4J: Failed to load class "org.slf4j.impl.StaticLoggerBinder".  
SLF4J: Defaulting to no-operation (NOP) logger implementation  
SLF4J: See http://www.slf4j.org/codes.html#StaticLoggerBinder for further details.
```

# View the Dataframe

```
>>> df.show()
```

```
hadoop@ip-172-31-43-113:~  
>>> df.show()  
+-----+-----+-----+-----+-----+-----+-----+-----+-----+  
| id| airline| date| location| rating| cabin| value| recommended| review|  
+-----+-----+-----+-----+-----+-----+-----+-----+-----+  
|10001|Delta Air Lines|21-Jun-14| Thailand| 7| Economy| 4| YES|Flew Mar 30 NRT t...|  
|10002|Delta Air Lines|19-Jun-14| USA| 0| Economy| 2| NO|Flight 2463 leavi...|  
|10003|Delta Air Lines|18-Jun-14| USA| 0| Economy| 1| NO|Delta Website fro...|  
|10004|Delta Air Lines|17-Jun-14| USA| 9| Business| 4| YES|"I just returned ...|  
|10005|Delta Air Lines|17-Jun-14| Ecuador| 7| Economy| 3| YES|"Round-trip fligh...|  
|10006|Delta Air Lines|17-Jun-14| USA| 9| Business| 5| YES|Narita - Bangkok ...|  
|10007|Delta Air Lines|14-Jun-14| UK| 0| Economy| 1| NO|Flight from NY La...|  
|10008|Delta Air Lines|14-Jun-14| USA| 0| Economy| 1| NO|Originally I had ...|  
|10009|Delta Air Lines|13-Jun-14| USA| 4| Business| 2| NO|We flew paid busi...|  
|10010|Delta Air Lines|13-Jun-14| UK| 9| Economy| 3| YES|"I flew from Heat...|  
|10011|Delta Air Lines|11-Jun-14| USA| 10| Economy| 4| YES|I was a bit stubb...|  
|10012|Delta Air Lines|10-Jun-14| Australia| 10| Economy| 5| YES|JFK-LHR. Had a gr...|  
|10013|Delta Air Lines| 9-Jun-14| USA| 0| Economy| 1| NO|My wife and I fly...|  
|10014|Delta Air Lines| 9-Jun-14| USA| 10| Premium| 5| YES|DL 1134 PBI-ATL. ...|  
|10015|Delta Air Lines| 6-Jun-14| USA| 0| Economy| 2| NO|Our flight from F...|  
|10016|Delta Air Lines| 5-Jun-14| USA| 0| Economy| 1| NO|On May 22 after a...|  
|10017|Delta Air Lines| 3-Jun-14| Canada| 9| Economy| 4| YES|Considering how D...|  
|10018|Delta Air Lines| 2-Jun-14| USA| 9| Economy| 5| YES|Travelled MSP-LHR...|  
|10019|Delta Air Lines|29-May-14| USA| 7| Business| 2| YES|JFK-LAX on a 757-...|  
|10020|Delta Air Lines|28-May-14| UK| 9| Economy| 3| YES|Third long haul f...|  
+-----+-----+-----+-----+-----+-----+-----+-----+-----+  
only showing top 20 rows  
>>> 
```



# Display descriptive stats using the describe function

```
>>> df.describe().show()
```

```
hadoop@ip-172-31-66-228:~
>>>
>>>
>>> df.describe().show()
```


summary	id	airline	date	location	rating	cabin	value	recommended	review
count	992	992	992	992	992	992	992	992	992
mean	10496.5	null	null	null	4.571572580645161	null	2.78125	null	null
stddev	286.5100347282796	null	null	null	3.803036434465363	null	1.5689762355914316	null	null
min	10001	American Airlines	1-Apr-14	Australia	0	Business	0	NO	"12 Jan flight 67...
max	10992	United Airlines	9-Sep-13	Venezuela	9	Premium	5	YES	with a group of 8...

```
>>>
>>>
>>>
```



Select only those observations where location is Ireland or Canada

```
>>> df[df.location.isin("Ireland","Canada")].show( )
```

 `hadoop@ip-172-31-66-228:~`

```
>>>
>>>
>>> df[df.location.isin("Ireland","Canada")].show( )
```

id	airline	date	location	rating	cabin	value	recommended	review
10017	Delta Air Lines	3-Jun-14	Canada	9	Economy	4	YES	Considering how D...
10025	Delta Air Lines	19-May-14	Canada	9	Economy	5	YES	YUL-JFK-NRT-HKG. ...
10026	Delta Air Lines	18-May-14	Canada	1	Economy	1	NO	Worst experience ...
10032	Delta Air Lines	13-May-14	Canada	10	Economy	5	YES	I flew from YYZ t...
10045	Delta Air Lines	1-May-14	Canada	2	Economy	2	NO	Flew Atlanta to S...
10047	Delta Air Lines	29-Apr-14	Canada	2	Economy	1	NO	The review is for...
10049	Delta Air Lines	28-Apr-14	Ireland	5	Economy	1	NO	Travelled Dublin ...
10058	Delta Air Lines	14-Apr-14	Canada	8	Economy	4	YES	I went DTW-SFO AN...
10070	Delta Air Lines	22-Mar-14	Canada	2	Economy	2	NO	Flew Jan 22 from ...
10078	Delta Air Lines	12-Mar-14	Canada	10	Economy	5	YES	Feb 25th flight D...
10093	Delta Air Lines	10-Feb-14	Canada	0	Economy	1	NO	"Wife and I flyin...
10094	Delta Air Lines	10-Feb-14	Ireland	7	Economy	3	YES	I flew from DUB-J...
10095	Delta Air Lines	10-Feb-14	Canada	4	Economy	4	NO	Flew JNB-ATL-SNA ...
10144	Delta Air Lines	29-Nov-13	Ireland	9	Economy	5	YES	Flew to New York ...
10166	Delta Air Lines	23-Oct-13	Canada	7	Economy	4	YES	"On our outbound ...
10213	Delta Air Lines	4-Sep-13	Ireland	10	Economy	5	YES	I recently flew w...
10243	Delta Air Lines	31-Jul-13	Canada	1	Economy	1	NO	07/19/2013 DL3655...
10251	US Airways	11-Jun-14	Canada	0	Economy	3	NO	US Airways now st...
10263	US Airways	25-May-14	Ireland	3	Economy	3	NO	"Delays on all bo...
10265	US Airways	22-May-14	Canada	2	Economy	3	NO	"Terrible. I had ...

only showing top 20 rows

Create a Spark temporary table (so that we can query the table using SparkSQL)

```
>>> df.createOrReplaceTempView("airlines")
```

```
>>> df.createOrReplaceTempView("airlines")
```

# Use SparkSQL to query the temporary table that you just loaded in the previous step

```
>>> sqlDF = spark.sql("SELECT * FROM airlines").show()
```

```
hadoop@ip-172-31-43-113:~  
>>> sqlDF = spark.sql("SELECT * FROM airlines").show()  
+-----+-----+-----+-----+-----+-----+-----+-----+-----+  
| id| airline| date| location| rating| cabin| value| recommended| review|  
+-----+-----+-----+-----+-----+-----+-----+-----+-----+  
|10001|Delta Air Lines|21-Jun-14| Thailand| 7| Economy| 4| YES|Flew Mar 30 NRT t...|  
|10002|Delta Air Lines|19-Jun-14| USA| 0| Economy| 2| NO|Flight 2463 leavi...|  
|10003|Delta Air Lines|18-Jun-14| USA| 0| Economy| 1| NO|Delta Website fro...|  
|10004|Delta Air Lines|17-Jun-14| USA| 9| Business| 4| YES|"I just returned ...|  
|10005|Delta Air Lines|17-Jun-14| Ecuador| 7| Economy| 3| YES|"Round-trip fligh...|  
|10006|Delta Air Lines|17-Jun-14| USA| 9| Business| 5| YES|Narita - Bangkok ...|  
|10007|Delta Air Lines|14-Jun-14| UK| 0| Economy| 1| NO|Flight from NY La...|  
|10008|Delta Air Lines|14-Jun-14| USA| 0| Economy| 1| NO|Originally I had ...|  
|10009|Delta Air Lines|13-Jun-14| USA| 4| Business| 2| NO|We flew paid busi...|  
|10010|Delta Air Lines|13-Jun-14| UK| 9| Economy| 3| YES|"I flew from Heat...|  
|10011|Delta Air Lines|11-Jun-14| USA| 10| Economy| 4| YES|I was a bit stubb...|  
|10012|Delta Air Lines|10-Jun-14| Australia| 10| Economy| 5| YES|JFK-LHR. Had a gr...|  
|10013|Delta Air Lines| 9-Jun-14| USA| 0| Economy| 1| NO|My wife and I fly...|  
|10014|Delta Air Lines| 9-Jun-14| USA| 10| Premium| 5| YES|DL 1134 PBI-ATL. ...|  
|10015|Delta Air Lines| 6-Jun-14| USA| 0| Economy| 2| NO|Our flight from F...|  
|10016|Delta Air Lines| 5-Jun-14| USA| 0| Economy| 1| NO|On May 22 after a...|  
|10017|Delta Air Lines| 3-Jun-14| Canada| 9| Economy| 4| YES|Considering how D...|  
|10018|Delta Air Lines| 2-Jun-14| USA| 9| Economy| 5| YES|Travelled MSP-LHR...|  
|10019|Delta Air Lines|29-May-14| USA| 7| Business| 2| YES|JFK-LAX on a 757-...|  
|10020|Delta Air Lines|28-May-14| UK| 9| Economy| 3| YES|Third long haul f...|  
+-----+-----+-----+-----+-----+-----+-----+-----+-----+  
only showing top 20 rows  
>>> 
```