



# Advanced Big Data: Distributed Databases

Dan Zaratsian

AI/ML Architect, Gaming Solutions @ Google

[d.zaratsian@gmail.com](mailto:d.zaratsian@gmail.com)

<https://github.com/zaratsian>

# TOPICS

- **Session 1: Course Intro, Trends, and Approach to AI/ML**
- **Session 2: SQL and NoSQL**
- **Session 3: Cloud Machine Learning Services**
- **Session 4: Distributed ML with Spark and Tensorflow**
- **Session 5: Cloud Generative AI Services and Architectures**
- **Session 6: Serverless ML, Architectures, and Deploying ML**



# Which Database should I use?

#GCPSketchnotes

@PVERGADIA THECLOUDGIRL.DEV

03.22.2023



RELATIONAL	RELATIONAL & NON-RELATIONAL	NON-RELATIONAL (NO SQL)	IN MEMORY	BARE METAL		
<b>Cloud SQL</b> Managed MySQL, PostgreSQL, SQL Server SLA: 99.95%	<b>AlloyDB</b> Managed PostgreSQL-compatible, with 100x faster analytics & 4x faster transaction queries SLA: 99.99%	<b>Spanner</b> Multi-dialect (PostgreSQL, GoogleSQL) database, unlimited scaling SLA: 99.999%	<b>Firestore</b> Serverless, document database with built-in cross-client sync and offline caching SLA: 99.999%	<b>Bigtable</b> Key-value store for large scale, low latency workloads, Supports HBase APIs SLA: 99.999%		
Fully managed experience for low-touch administration with improved availability, security & governance, support						
Good For:						
- Fastest lift & shift migrations - OLTP workloads - Managed experience - Common API & control plane across database engines	- PostgreSQL workloads that need high performance & high availability or scale - HTAP (hybrid transaction/analytical processing) - Migrations off commercial databases	Highest scale and availability requirements without compromising on SQL capabilities, consolidating multiple databases for cost savings	- Rapid and cost-efficient application development (no need for middle tier) - Easy aggregation from multiple data sources	- Cost-sensitive applications - Need high throughput & consistent single-digit millisecond latencies irrespective of scale		
Use Case:						
Web Frameworks ERP CRM Ecommerce & Web	Operational Analytics ERP, CRM Financial Services Ecommerce & Web SaaS	Order & Inventory Management Online Banking, payments & ledger Gaming: player profiles & gameplay data Electronic Medical Records Catalog Metadata Management	Web & Mobile Apps News Feeds, Social Chat, Influencer Engagements Game Saves, Player Profiles Retail Catalogs, Point of Sales	Real-Time Analytics (Personalization, Fraud detection) IoT/Clickstream/Time Series Feature Stores, Operational Data Hubs & Data Fabrics Batch Unstructured Data Processing Financial Markets, Crypto Ledgers	Database caching Session Store Jobs and Queues Leaderboards Fast Data Ingestion	Legacy Applications Datacenter Exits

# Google BigQuery

Proprietary + Confidential

Google Cloud Platform's  
**enterprise data warehouse**  
for analytics

Familiar **SQL 2011** query  
language and functions

**Exabyte-scale, serverless**  
data warehousing

**Encrypted, durable, secure,**  
And highly available



Built-in support for Generative AI

Unique

**Real-time** insights from streaming data

Unique

Built-in **ML and Geospatial** for  
predictive insights

Unique

High-speed, in-memory **BI Engine**  
for faster reporting and analysis

Unique

Nested and repeated fields,  
user-defined functions in JavaScript

Unique

# Your database choice depends on your needs



Availability



Multi-Region



Skill Set



Operations



Compatibility



Consistency



Cost



Future proofing



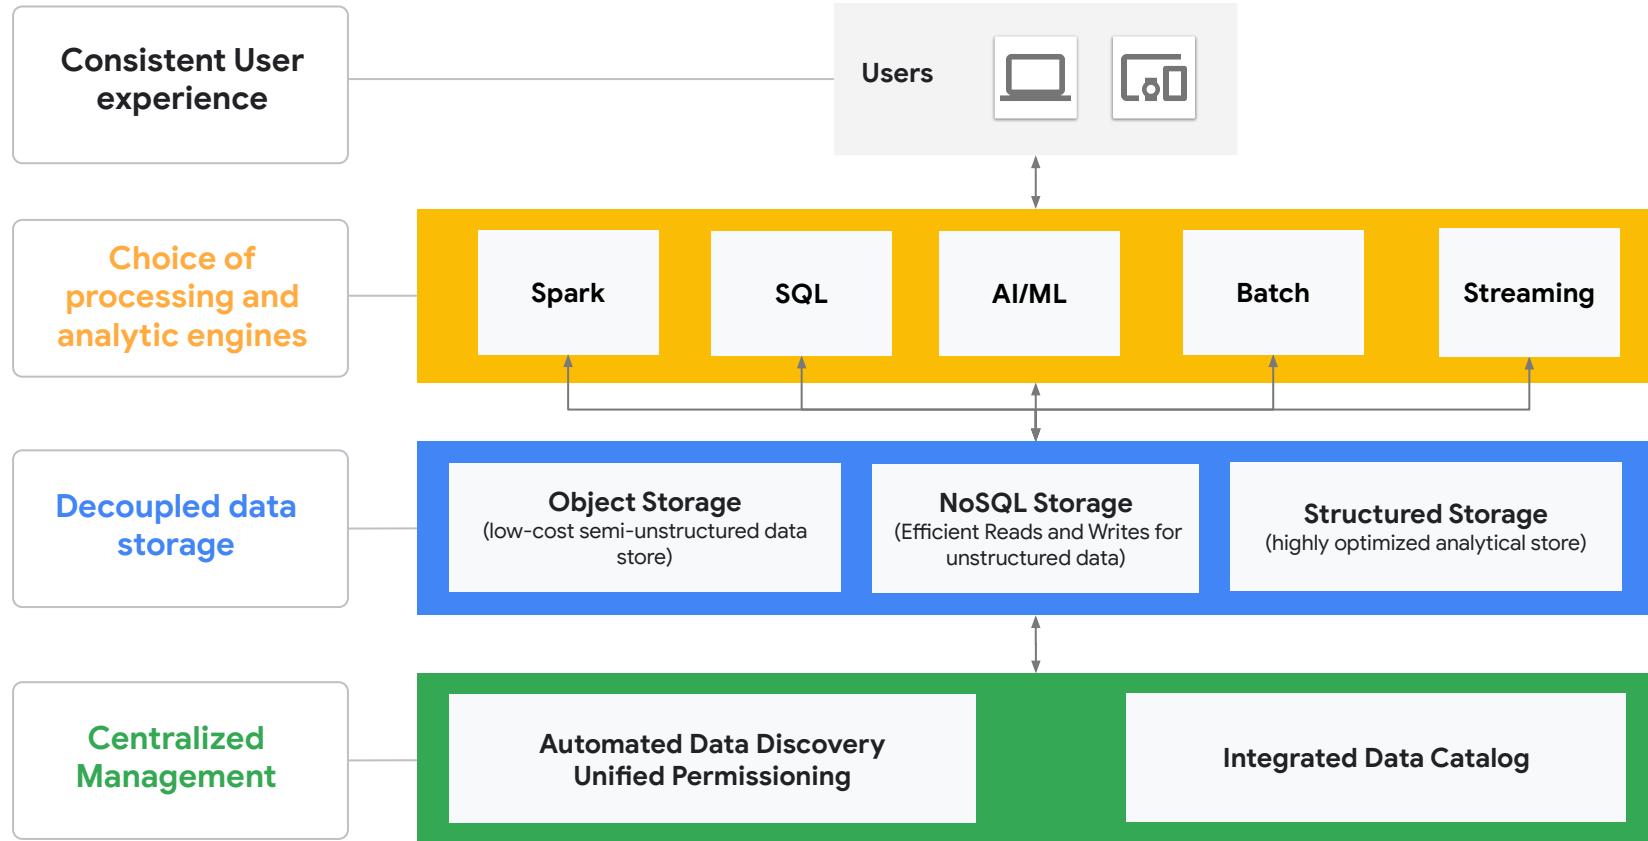
Data model



Open standard

# Data Warehouse building blocks

Proprietary + Confidential



# SQL vs NoSQL

## SQL

[MySQL, Postgres, BigQuery, Oracle, etc.]

- Relational databases
- Data stored in tables (tabular)
- Requires schema
- Mostly vertically scalable
- Generally better for multi-row transactions
- Best for heavy transactional workloads
- Uses SQL standard languages for queries

## NoSQL

[HBase, Cassandra, MongoDB, Firestore, etc.]

- Unstructured format
- Non-relational database
- Does **not** require a schema
- Easily scalable and distributed
- Horizontally scalable
- Low-latency scans and lookups
- Best for cases where query path is well defined
- Query language not as standardized and syntax is specific to DB.

Note: New DBs are blurring the line. SQL databases now support NoSQL and nested JSON, while NoSQL databases allow users to write SQL queries.

# SQL



Google  
BigQuery



Amazon Athena



Amazon **Redshift**



Azure  
Synapse  
Analytics

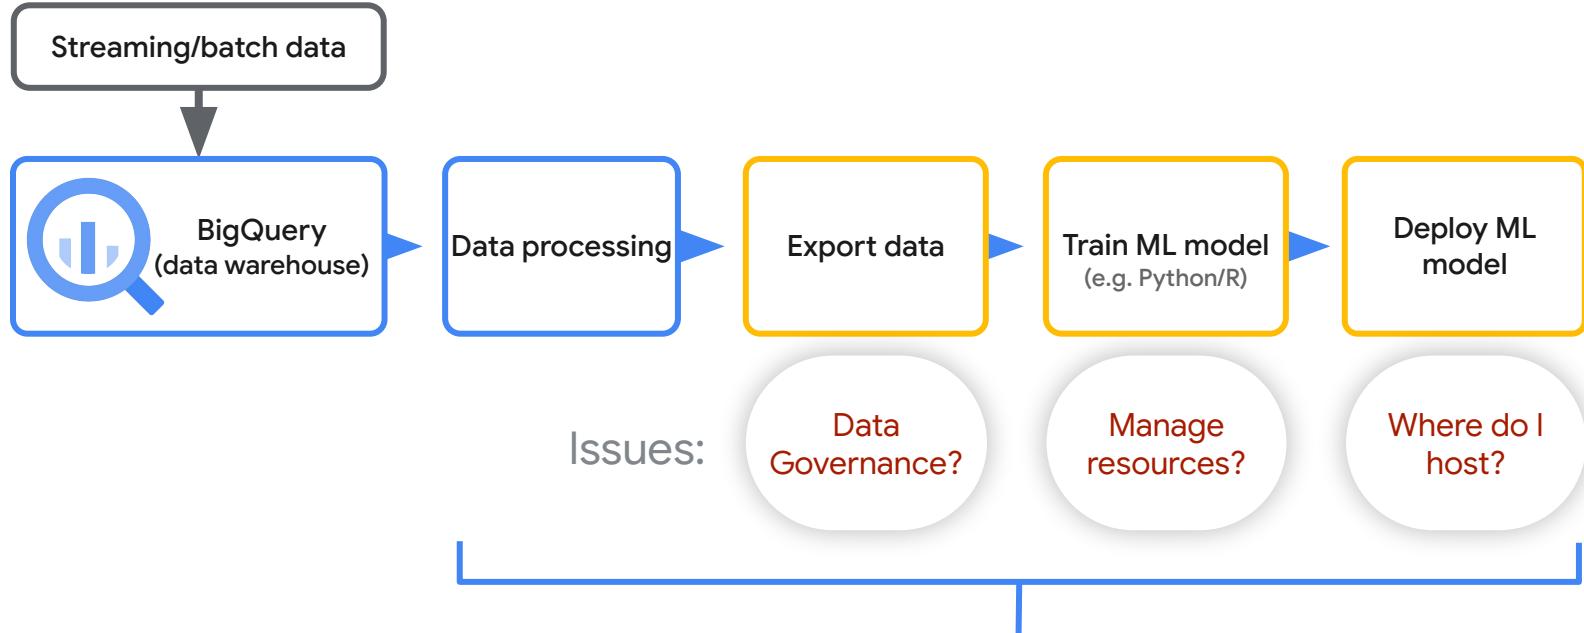




# BigQuery

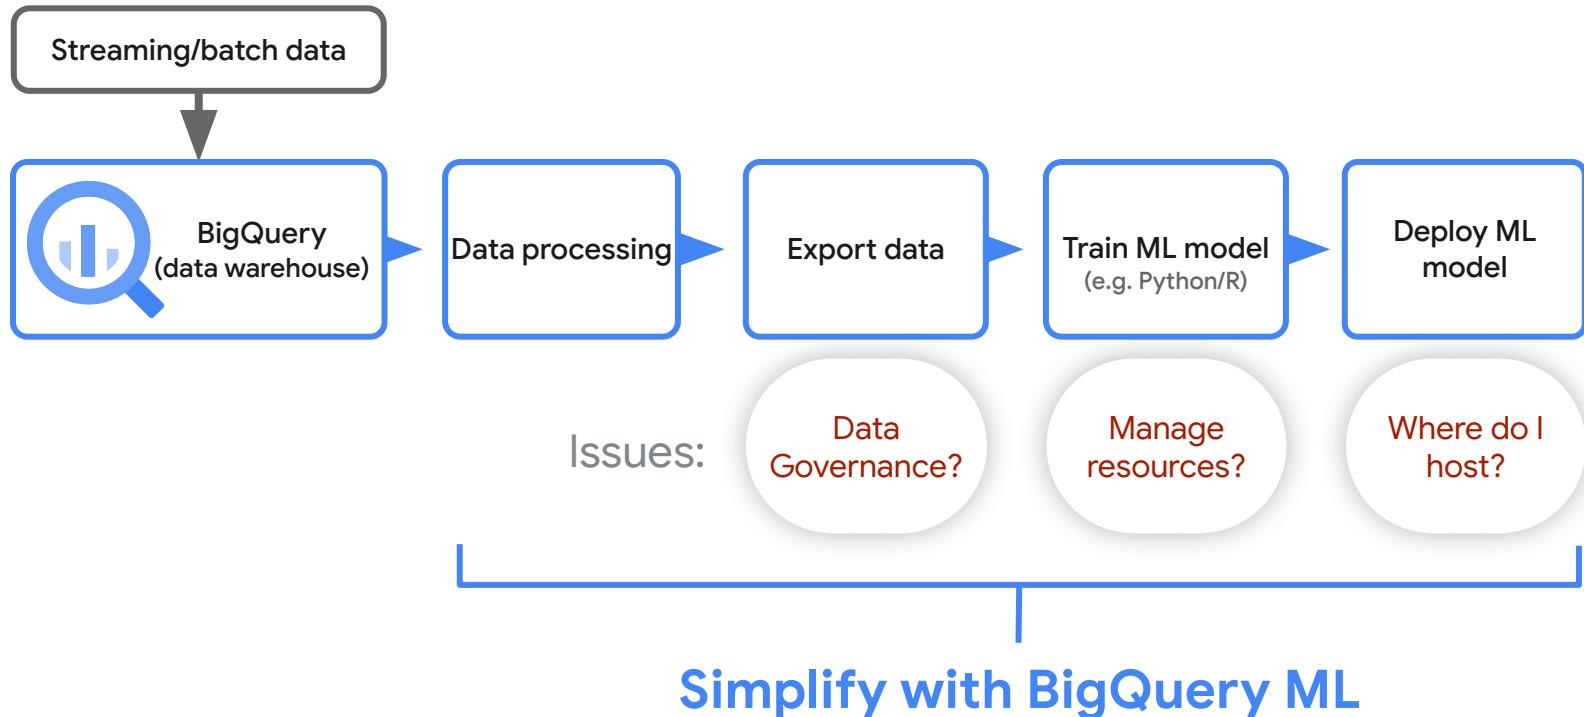
Serverless and cost-effective enterprise data warehouse that works across clouds and scales with your data

# Typical ML workflow poses many challenges



Multiple products & roles can lead to unnecessary complexity & costs

# BigQuery ML greatly simplifies the ML workflow



# BigQuery ML brings BigQuery and Vertex AI together



# BigQuery ML integrates with Vertex AI in many areas

## Model Training

- Vertex Training for training Tensorflow and XGBoost models
- Vertex Tables for training AutoML tabular models  
AutoML and Forecasting On roadmap models
- Vertex Vizer for hyperparameter tuning

## Pre-trained Models

- Inference with LLM foundation models
- Inference for images with Vision API
- Inference for text with NLP API
- Inference for text with Translate API
- Inference for docs with Document AI On roadmap

## MLOps

- Register models with Vertex AI Model Registry
- Export models and deploy them to Vertex AI Prediction and Vertex Explainable AI
- Model monitoring with Vertex AI
- Vertex AI Pipeline components for use in Notebooks

# ML with just a few SQL statements

## 1 Train



```
CREATE OR REPLACE MODEL `bqml.penguins_model`  
OPTIONS (model_type='linear_reg',  
        input_label_cols=['body_mass_g']) AS  
SELECT * FROM `public-data.ml_datasets.penguins`  
WHERE body_mass_g IS NOT NULL
```

## 2 Predict



```
SELECT * FROM  
ML.PREDICT (  
    MODEL `bqmlTutorial.penguins_model`,  
    (SELECT * FROM `public-data.ml_datasets.penguins`  
    WHERE body_mass_g IS NOT NULL AND  
          island = "Biscoe"))
```

**Train and deploy** ML models in SQL

**Execute** ML workflows without moving data from BigQuery

**Built-in** infrastructure management, security & compliance

# BQML features & capabilities

## Rich Built-in Models

- **Supervised Learning:** Linear / Logistic regression, Boosted trees, Random Forest, DNN, Wide & Deep models, AutoML Tables
- **Unsupervised Learning:** K-Means, PCA, Matrix Factorization, Autoencoder
- **Time Series:** ARIMA\_PLUS (univariate) and ARIMA\_PLUS\_XREG (multivariate)

## Feature Engineering & others

- Rich feature engineering support
- Model evaluation
- Model management

## ML Ops & Cycles

- Hyperparameter Tuning
- Explainable AI
- Integration with Vertex Model Registry
- Components for Vertex Pipelines
- Model Monitoring (soon)

## AI/ML Applications

- Generative AI
- NLP & Translation
- Computer Vision
- Embeddings
- Recommendation
- Forecasting
- Anomaly Detection
- Speech & Document AI

## Inference Engine

- **Remote Model Inference:** with Vertex Endpoint
- **Imported Model Inference:** with TF, TFLite, XGBoost & ONNX (enabling sklearn, Pytorch, SparkML)
- **Export Modes:** for online inference in Vertex



# Predict against Vertex LLMs in BigQuery

```
SELECT * FROM
ML.GENERATE_TEXT (
  MODEL `my_company.llm_model`,
  (SELECT
    CONCAT("Give the country name for city: ", city) AS prompt
    FROM example_table),
  STRUCT (0.2 AS temperature,
          1024 AS max_output_tokens,
          0.8 AS top_p,
          40 AS top_k))
```



---

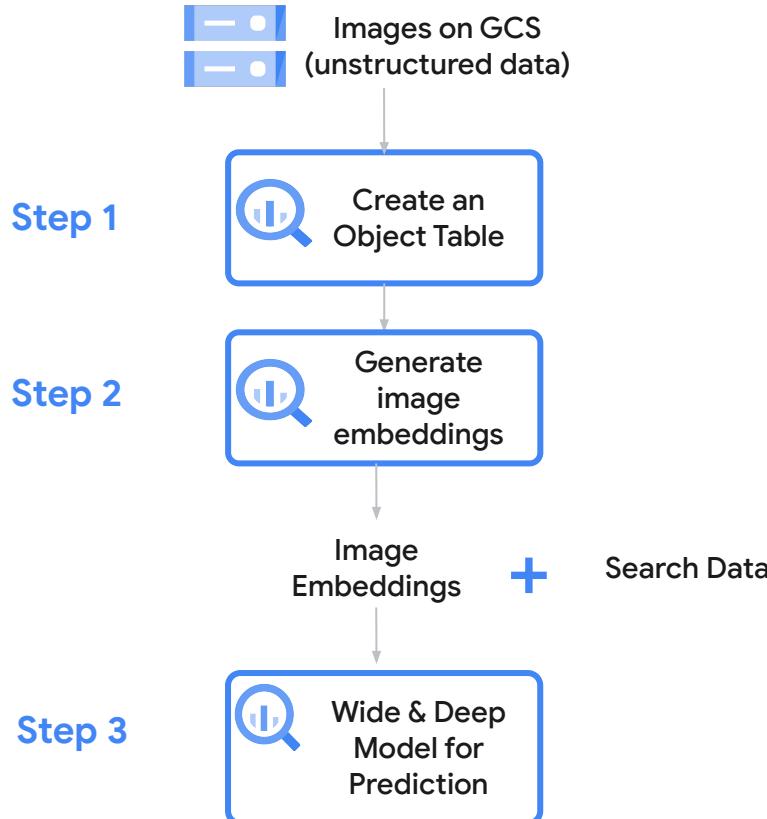
**Use LLM models in SQL**

---

**Execute Generative AI  
workflows inside BigQuery**

# BigQuery Example (SQL + NoSQL) Data processing

Predict Click-through rate of rentals (based on image and structured data)



Now you can combine  
structured and unstructured  
data in BQML

# Access to image data by creating an Object Table



```
CREATE OR REPLACE EXTERNAL TABLE
`images.property-images`
WITH CONNECTION `us.property-images`
OPTIONS(
object_metadata="DIRECTORY",
uris=["gs://demo/images/*"]
);
```



uri	content_type	size
gs://as-bqml-launch-demo/images/ER010-aerialrearext.jpg	image/jpeg	1940969
gs://as-bqml-launch-demo/images/KD1111-rearext.jpg	image/jpeg	1508587
gs://as-bqml-launch-demo/images/ER011-aerialrearext.jpg	image/jpeg	1949278
gs://as-bqml-launch-demo/images/B800-rearext.jpg	image/jpeg	338409
gs://as-bqml-launch-demo/images/b372-rearext.jpg	image/jpeg	1841978
gs://as-bqml-launch-demo/images/J10975-aerial-2.jpg	image/jpeg	1606464

Step 1: Create an object table

Step 2: Generate image embeddings

Step 3: Wide & Deep model for prediction

# Create image embeddings by importing a Tensorflow image model



```
CREATE OR REPLACE MODEL  
`pipeline.resnet_imagenet_embeddings_model`  
OPTIONS(  
  model_type="TENSORFLOW",  
  color_space="RGB",  
  model_path="gs://demo/resnet-embeddings/*")
```



```
SELECT *  
FROM ML.PREDICT (  
  MODEL `pipeline.resnet_imagenet_embeddings_model`,  
  (SELECT * FROM `images.property-images`))
```



img_name	pc1	pc2	pc3
EC1-Bar.jpg	1.8493...	1.3984...	1.9969...
EC1-ext.jpg	-18.16...	-8.0527...	1.3630...
EC1-Bar2.jpg	-1.382...	-0.5212...	-3.135...
EC1-Bar3.jpg	3.8091...	-5.0284...	2.1694...
ec1-pool.jpg	-4.891...	7.1742...	-5.182...

Step 1: Create an object table

Step 2: Generate image embeddings

Step 3: Wide & Deep model for prediction

# Train a Wide and Deep model for prediction

image\_embeddings

img_name	pc1	pc2	pc3
EC1-Bar.jpg	1.8493...	1.3984...	1.9969...
EC1-ext.jpg	-18.16...	-8.0527...	1.3630...
EC1-Bar2.jpg	-1.382...	-0.5212...	-3.135...
EC1-Bar3.jpg	3.8091...	-5.0284...	2.1694...
ec1-pool.jpg	-4.891...	7.1742...	-5.182...

search\_data

img_name	distanceToB...	location	state
EC1-Bar.jpg	Oceanfront	Corolla	NC
EC1-ext.jpg	Oceanfront	Corolla	NC
EC1-Bar2.jpg	Oceanfront	Corolla	NC
EC1-Bar3.jpg	Oceanfront	Corolla	NC
ec1-pool.jpg	Oceanfront	Corolla	NC

img_name	distanceTo...	location	state	pc1	pc2	pc3
EC1-Bar.jpg	Oceanfront	Corolla	NC	1.84936261...	1.39845214...	1.99695636...
EC1-ext.jpg	Oceanfront	Corolla	NC	-18.166918...	-8.0527233...	1.36307251...
EC1-Bar2.jpg	Oceanfront	Corolla	NC	-1.3826520...	-0.5212412...	-3.1352281...
EC1-Bar3.jpg	Oceanfront	Corolla	NC	3.80916701...	-5.0284545...	2.16948761...
ec1-pool.jpg	Oceanfront	Corolla	NC	-4.8918667...	7.17424064...	-5.1822880...

image\_embeddings\_with\_search

Step 1: Create an object table

Step 2: Generate image embeddings

Step 3: Wide & Deep model for prediction



```
CREATE OR REPLACE MODEL pipeline.wide_and_deep
OPTIONS(MODEL_TYPE='DNN_LINEAR_COMBINED_CLASSIFIER'
      ,
      AUTO_CLASS_WEIGHTS = TRUE,
      ACTIVATION_FN = 'RELU',
      BATCH_SIZE = 64,
      DROPOUT = 0.1,
      EARLY_STOP = FALSE,
      HIDDEN_UNITS = [128, 64, 32, 64, 128],
      INPUT_LABEL_COLS = ['clicked'],
      LEARN_RATE=0.00001,
      MAX_ITERATIONS = 25,
      OPTIMIZER = 'RMSPROP',
      DATA_SPLIT_METHOD="CUSTOM",
      DATA_SPLIT_COL = "split_col_bool",
      ENABLE_GLOBAL_EXPLAIN = TRUE
    )
```

Select \* from image\_embeddings\_with\_search

Google Cloud

# Use ML.predict to predict the click-through rate of rental listings



```
SELECT  
  *  
FROM  
  ML.PREDICT(MODEL  
    `bqml-demo.pipeline.wide_and_deep`,  
    (SELECT * FROM base));
```



name	img_name	avg_prob
Aqua Dream	J10975-aerial-2.jpg	0.77494834...
Atlantis	KD1111-rearext.jpg	0.78830668...
Creme de la Creme	J20944-aerialrearext.jpg	0.69067030...
Heavens to Betsy	ER011-aerialrearext.jpg	0.82573245...
Mariner's Compass	ER008-aerialrearext.jpg	0.74198232...
SEAesta	KD1301-rearext.jpg	0.69142866...
Salty Paws	B800-rearext.jpg	0.70336384...
Sound to Sea Beach Club	ec3-aerialrearext.jpg	0.78168062...

Step 1: Create an object table

Step 2: Generate image embeddings

Step 3: Wide & Deep model for prediction

# BigQuery Best Practices

- Don't SELECT \* unless you need every field
- Filter early and often using WHERE clauses
- Do the biggest joins first
- Avoid self-join if you can, since it squares the number of rows processed
- Built-in functions are faster than User Defined Functions (UDFs)

# BigQuery Public Datasets



NYC OpenData

citibike





# Cloud SQL

Managed Database Platform for MySQL, PostgreSQL and SQL Server Databases

# Cloud SQL

Fully managed relational database service



**Supports PostgreSQL, MySQL and SQL Server**

Full compatibility with source database engines

**Fully Managed & Enterprise Ready**

Easy to set up, operate, and scale

**Trusted**

Enterprise-grade data protection, security and governance

**Developer Friendly**

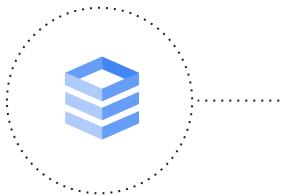
Application centric observability and API-first administration

More than

**90%**

of Google Cloud's top 100 customers use Cloud SQL

# When to choose Cloud SQL



-  Full compatibility with MySQL, PostgreSQL and SQL Server databases
-  Built-in high availability along with option for cross-region read replication
-  Automate database provisioning, storage capacity management, and other time-consuming tasks
-  Industry leading observability for PostgreSQL and MySQL with Query and System Insights
-  Proactive database wellness with Active Assist feature
-  Lift and shift migration of databases from on-prem or other clouds
-  Common control plane and API across MySQL, PostgreSQL and SQL Server
-  Integration with analytics and data visualization with BigQuery and Looker



## Considerations



Zero downtime for maintenance



Global database for read/write



Super-Admin Access to tune features

## With Cloud SQL you can

Innovate faster    Reduce risk    Gain observability



# Spanner

Enterprise-grade, globally-distributed, and strongly-consistent managed SQL database service built for the cloud

# What is Spanner?



## Relational

ACID transactions,  
SQL, Schemas



## Horizontally scalable

Distributed RDBMS,  
Near unlimited scale



## Fully managed ++

Simplified administration,  
Enterprise grade



99.999% uptime SLA



No maintenance downtime



Online, unlimited scaling



Automatic sharding



Automatic failure recovery  
RPO = 0, RTO = 0



Strong external consistency



Superior price-performance



Zero-touch global replication



Security and compliance

Spanner processes over **2 billion requests per second** at peak  
Spanner has more than **6 exabytes of data** under management

# When to choose Spanner



- Fully managed relational database with unlimited horizontal scale
- Need high availability with zero downtime and online schema changes
- High performance ACID transactions with strong consistency over regions
- PostgreSQL interface to empower developers
- Unified analytics and AI on transactional data with BigQuery and Vertex AI
- In-house real-time CDC and replication with Spanner change streams



## Considerations



You need low-latency local writes across regions



Your schema is unknown or highly sparse



You need a drop-in replacement for your legacy database

## With Spanner you get

Zero Downtime   Unlimited scale   99.999% uptime

# NoSQL

# Types of NoSQL

## Key-Value Stores



Cloud Firestore



**redis**



Amazon DynamoDB

## Column Oriented



Cloud Bigtable



*cassandra*

## Document DBs



Amazon Neptune



Titan



# Bigtable

Globally distributed, fully managed NoSQL database with high performance  
at any scale



# Bigtable

Real-time data serving and operational analytics at any scale

## Flexibility at scale

Flexible schema, eventual consistency\*

## High throughput

Millions of RPS, Predictable single-digit ms latency

## Compatible

HBase API, Apache Spark, Integrates with Apache Beam ecosystem

## Example use cases



Personalization



Time Series



Fraud detection



Product/Content metadata



Data Fabric/Operational Data Store



Customer 360

Bigtable has over **10 Exabytes of data** under management

Bigtable processes more than **5 Billion** requests per second at peak

\* Strong consistency within a single cluster

Battle tested by Google



Google Cloud

# When to choose Bigtable



- Build responsive applications with consistent single-digit ms latency
- NoSQL database to seamlessly scale to match storage and throughput needs
- Greater flexibility to evolve your applications, or need for sparse schemas
- Ensure high availability with multi-primary replication in up to 8 regions
- Modernize Apache HBase database with no downtime
- Unified analytics and AI on transactional data with BigQuery integration
- Use cases such as personalization, fraud detection, real-time analytics, IoT and more



## Considerations



Full SQL support



Need secondary indexes



Need drop-in replacement for Cassandra or DynamoDB

## With Bigtable you get

Autoscaling

Low Latency

99.999% uptime



# Firestore

Serverless NoSQL document database built for automatic scaling, high performance in the cloud

# Firestore

Unlock application innovation with [simplicity](#), [speed](#) and [confidence](#)



## Serverless, document database

JSON-compatible data model, serializable ACID transactions, elastic scalability, up to 99.999% availability SLA, pay only for what you consume



## Secure, backend as a service

Connect directly and securely to the database, making middle tiers optional



## Real-time sync & offline access

Built-in data syncing, and fallback to on-device caching when a client loses network connectivity



## Well integrated

Deliver results faster with native integrations with Google Cloud, Firebase and 3rd party developer services via Extensions



## Firestore by the numbers

Over

**4 million**

databases have been created in Firestore

Firebase apps power more than

**1 billion**

monthly active end-users using Firebase Auth

# When to choose Firestore



- Document database that effortlessly scales to meet any demand, with no maintenance, and with ACID support
- Accelerate development of mobile, web and IoT apps with direct connectivity to the database
- Built-in live sync and offline mode to make it easy to develop real-time applications
- Libraries for popular languages like Node.js, Java, Go, Ruby and PHP
- Seamless integration with Firebase



## Considerations



You need capacity-based pricing



Your data is relational in nature



You need customer-managed encryption keys

## With Firestore you get

Serverless      Unlimited scale      99.999% uptime



# Memorystore

Fully Managed **Redis** and **Memcached** database for sub-millisecond data access

# For workloads that require microseconds latency, Memorystore is a scalable, secure, highly available in-memory service

## Example use cases



Caching



Leaderboard



Real-time analytics



Jobs and queues



Session store



Fast data ingestion



Fully compatible with Redis and Memcached, offering easy migration for Redis and Memcached workloads.

Over **90%** of the top 100 Google Cloud customers use Memorystore.

# When to choose Memrystore



- Need fully managed Redis or Memcached database service
- Build application caches that provide sub-millisecond data access
- Scalability when application grows, and high availability (99.9%)
- Multiple read replicas with up to 6 nodes to serve read traffic
- Web, Mobile Application, Gaming Leaderboard
- Social, chat, news feed applications



## Considerations



Higher data persistence needed



Lack of SQL Support



Transaction support is not full  
ACID compliant

## With Memrystore, you get

Scale

High Availability

Choice of engine

# Assignment

# SQL & NoSQL Assignments

SQL Assignment - Please complete this as an individual assignment and upload your code as a notebook file, script, google doc, etc.

NoSQL Assignment - Please complete this as an individual assignment and upload your code as a notebook file, script, google doc, etc.

**Due Date:** Thursday, March 14 at 11:59pm EST

Email with any questions: [d.zaratsian@gmail.com](mailto:d.zaratsian@gmail.com)

# **Google BigQuery Sandbox Environment**

<https://console.cloud.google.com/bigquery>





Colaboratory, or "Colab" for short, allows you to write and execute Python in your browser, with

- Zero configuration required
- Free access to GPUs
- Easy sharing

<https://colab.sandbox.google.com/>

+ Code + Text

```
[13] from pyspark.ml.feature import VectorIndexer, VectorAssembler, StringIndexer  
[13] from pyspark.ml.evaluation import RegressionEvaluator  
[13] from pyspark.ml import Pipeline  
[13] from pyspark.ml.evaluation import MulticlassClassificationEvaluator
```

```
[9] spark = SparkSession.builder.appName("Bikesharing SpaksSQL").master("local[*]").getOrCreate()
```

```
[14] bikeshare_stations = spark.read.load('bikeshare_stations.csv', format="csv", header=True, inferSchema=True)
```

```
[16] bikeshare_stations.show(5)
```

```
▶ +-----+-----+-----+-----+-----+  
|station_id|      name|status|latitude|longitude|      location|  
+-----+-----+-----+-----+-----+  
|     3793| Rio Grande & 28th|active|30.29333|-97.74412|(30.29333, -97.74...|  
|     3291| 11th & San Jacinto|active|30.27193|-97.73854|(30.27193, -97.73...|  
|     4058| Hollow Creek & Ba...|active|30.26139|-97.77234|(30.26139, -97.77...|  
|     3797| 21st & University|active|30.28354|-97.73953|(30.28354, -97.73...|  
|     3838| Nueces & 26th|active|30.29068|-97.74292|(30.29068, -97.74...|  
+-----+-----+-----+-----+-----+  
only showing top 5 rows
```

```
▶ play bikeshare_stations.createOrReplaceTempView("bikeshare_stations")  
spark.sql("SELECT name, count(*) as count FROM bikeshare_stations group by name order by count desc").show(truncate=False)
```

```
▶ +-----+-----+  
|name|count|  
+-----+-----+  
|Lavaca & 6th| 2 |  
|Lake Austin & Enfield| 1 |  
|ACC - West & 12th Street| 1 |  
|Convention Center / 3rd & Trinity| 1 |
```

# Thanks!

# Advanced Big Data

**Dan Zaratsian**

AI/ML Architect, Gaming Solutions @ Google

[d.zaratsian@gmail.com](mailto:d.zaratsian@gmail.com)

<https://github.com/zaratsian>

# Next up...

## Session 2 SQL & NoSQL



## Session 3 Distributed ML



## Session 4 GenAI



GCP PubSub

## Session 5 Serverless



Cloud Functions



Cloud Run



Vertex AI

## Session 6 Cloud



Vertex AI