

# Natural Language Processing Workshop

**Amy Hemmeter, MSA Class of '18**  
**Manager/Data Scientist/Machine Learning Engineer**

“We are seeing a surge in demand for NLP applications across sectors and use cases. Our clients are implementing NLP solutions for deeper customer insights, risk mitigation, and operational efficiency.”

- Nathan Peifer, formerly of EY, now at Elicit, personal communication, 2020

“We’re looking for better skills for dealing with unstructured data, NLP is becoming increasingly important as much of the data available is text.”

- Garima Sharma, Manager, Enterprise Data Intelligence at Domino’s (Women in Data Science Conference 2020, Detroit)

“NLP is full of fraudsters”

- Amy

# Natural Language Processing

- NLP is a subfield of artificial intelligence that deals with understanding and in some cases producing, human (“natural”) language
- We’re going to cover the following NLP topics:
  - Word vectors
  - Deep learning solutions in NLP
  - Transformers
    - BERT
    - GPT
  - High-level descriptions of NLP tasks

# Rule-Based NLP

- A series of rules in code that very explicitly spelled out how a computer is supposed to understand human language
- Often based on keywords and collocations (words that occur together)
- Regular expressions play a big role
- Very time-consuming and expensive for the company
- Fairly old-fashioned, no one's first choice and doesn't show up on job postings that request NLP skills very often
- However, many systems are still partially rule-based -- including those of many MSA employers (from my experience of talking to them in 2018)

# Expert AI Revealed To Be 1,000,000 If-Else Statements Stacked in a Trenchcoat



“I thought the movie was terribly well done, bravo to all the actors.”

“Garbage”.



1. Words are not numbers

1. Words are not numbers
2. Input can be different lengths

1. Words are not numbers
2. Input can be different lengths
3. Words can mean different things in different contexts

**1. Words are not numbers**

2. Input can be different lengths

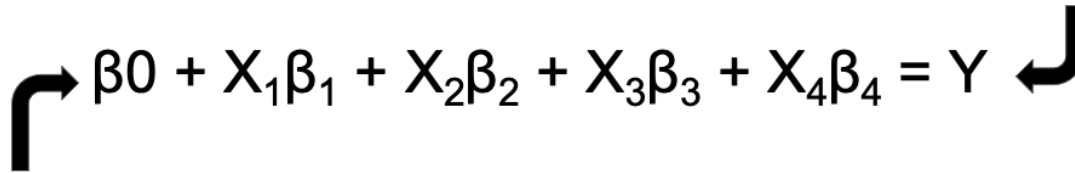
3. Words can mean different things in different contexts

# Word Embeddings

## Traditional Dataset

Number of Bedrooms	Number of Bathrooms	Age in Years	Walkability
2	1	10	87.2
5	3	2	50.2
3	1.5	25	95.2

**Output  
is a  
number**


$$\beta_0 + X_1\beta_1 + X_2\beta_2 + X_3\beta_3 + X_4\beta_4 = Y$$

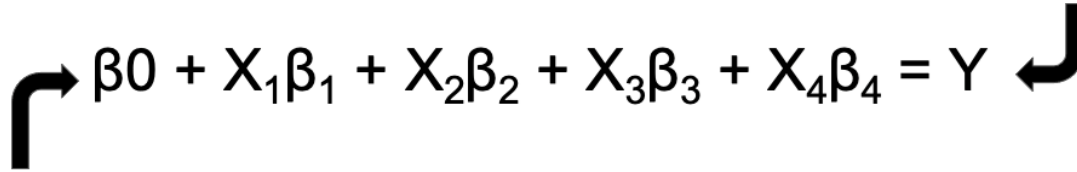
**Inputs are numbers**

# Word Embeddings

## Traditional Dataset

Number of Bedrooms	Number of Bathrooms	Age in Years	Walkability
2	1	10	87.2
5	3	2	50.2
3	1.5	25	95.2

**Output  
is a  
number**


$$\beta_0 + X_1\beta_1 + X_2\beta_2 + X_3\beta_3 + X_4\beta_4 = Y$$

**Inputs are numbers**

Input to ML model is a vector: [2 1 10 87.2]

# Word Vectors

- How do we turn a word into a number?
- We could use “one-hot” vectors - each vector is the length of your vocabulary, and a 1 denotes your word in that vocabulary

Vocabulary:  
Man, woman, boy,  
girl, prince,  
princess, queen,  
king, monarch



	1	2	3	4	5	6	7	8	9
man	1	0	0	0	0	0	0	0	0
woman	0	1	0	0	0	0	0	0	0
boy	0	0	1	0	0	0	0	0	0
girl	0	0	0	1	0	0	0	0	0
prince	0	0	0	0	1	0	0	0	0
princess	0	0	0	0	0	1	0	0	0
queen	0	0	0	0	0	0	1	0	0
king	0	0	0	0	0	0	0	1	0
monarch	0	0	0	0	0	0	0	0	1

Each word gets  
a 1x9 vector  
representation

[Source](#)

# Problems with One-Hot

- Real vocabularies can be millions of words, meaning your vectors could be incredibly high-dimensional
- Very clunky to work with
- No information encoded in the word vector
- All vectors are orthogonal, no one is more similar to the other than any of the others
- What happens if a word occurs that wasn't in your training set?



# Distributional Semantics

“You shall know a word by the company it keeps” - J.R. Firth

- We draw on a linguistic notion called “distributional semantics”
- We can get a large number of contexts for a word and use it to build a vector representation of that word

*...government debt problems turning into **banking** crises as happened in 2009...*

*...saying that Europe needs unified **banking** regulation to replace the hodgepodge...*

*...India has just given its **banking** system a shot in the arm...*

## Quick quiz:

Does this sound like anything you've learned about in your text analytics class?  
Defining word vectors by their context?

# Word Embeddings

	$X_1$	$X_2$	$X_3$	$X_4$	$X_5$
pasta	0.02182935	0.72359903	0.35498939	0.79730453	0.76512217
hurtled	0.53065205	0.27409522	0.95427668	0.75399903	0.4420345
the	0.47452104	0.18728204	0.26258075	0.90970631	0.62042977
said	0.559265	0.39511139	0.02820151	0.03384123	0.30565618
with	0.48774497	0.63185696	0.34809497	0.01427407	0.47724118

# Word Embeddings

	$X_1$	$X_2$	$X_3$	$X_4$	$X_5$
pasta	0.02182935	0.72359903	0.35498939	0.79730453	0.76512217
hurtled	0.53065205	0.27409522	0.95427668	0.75399903	0.4420345
the	0.47452104	0.18728204	0.26258075	0.90970631	0.62042977
said	0.559265	0.39511139	0.02820151	0.03384123	0.30565618
with	0.48774497	0.63185696	0.34809497	0.01427407	0.47724118

Input to ML model is a vector: [0.02182935 0.72359903 0.35498939 0.79730453 0.76512217]

# Word meaning visualized



Source: Stanford cs224 notes

# Word meaning visualized




Source: Stanford cs224 notes

# Word Vector Arithmetic




Source: [the morning paper](#)

# Word Vector Arithmetic





# Word Vector Arithmetic



# How do we get these vectors?

- Training a neural network with one-hot vector encodings as inputs to do a “fake” Language Modeling task
- We then pop out the hidden layer of that neural network and use that as our embedding

# What is language modeling?

- The task of predicting the probability of a sentence

“I’ll text you when I get \_\_\_\_\_”

# Other Uses for Language Modeling

- Speech recognition
  - We want to know that the sounds that are very similar in “wreck a nice beach” are more likely “recognize speech” -- we want to get a probability of the sentence
- Machine Translation
  - Change idiomatic “large winds tonight” from another language to more naturally English “High winds tonight”
- Spelling correction
  - “The office is about 15 minuets from my house” → “the office is about 15 minutes from my house”
- Anytime you need to know the probability of a sentence!


## Quick quiz:

In some forms of text analytics, you remove so-called “stop words”  
- function words like “a”, “the”, “of”, “in”, “and”, etc.

Do you think you should remove stop words in a language modeling task? Why or why not?

Answer: NO.

“I want go \_\_\_\_\_”



“I want to go to the  
\_\_\_\_\_”

“I want to go \_\_\_\_\_”

# word2vec

- Word2vec (Mikolov et al. 2013) is a framework for learning word vectors
- Inputs are vectors the same length as your vocabulary - so if you have 10,000 words you have a vector of dimension 10,000
- There are two versions of word2vec
  - SkipGram - predict context words from center word
  - Continuous Bag of Words - predict center word from (bag of) context words

## 2. Sliding Window

derekchia.com

#1	natural Xκ	language Y(c=1)	processing Y(c=2)	and	machine	learning	is	fun	and	exciting	#1
#2	natural	language Y(c=1)	processing Xκ	and Y(c=2)	machine	learning	is	fun	and	exciting	#2
#3	natural	language Y(c=1)	processing Y(c=2)	and Xκ	machine Y(c=3)	learning	is	fun	and	exciting	#3
#4	natural	language Y(c=1)	processing Y(c=2)	and Xκ	machine Y(c=3)	learning	is	fun	and	exciting	#4
#5	natural	language Y(c=1)	processing Y(c=2)	and Xκ	machine Y(c=3)	learning	is	fun	and	exciting	#5
#6	natural	language Y(c=1)	processing Y(c=2)	and Xκ	machine Y(c=3)	learning	is	fun	and	exciting	#6
#7	natural	language Y(c=1)	processing Y(c=2)	and Xκ	machine Y(c=3)	learning	is	fun	and	exciting	#7
#8	natural	language Y(c=1)	processing Y(c=2)	and Xκ	machine Y(c=3)	learning	is	fun	and	exciting	#8
#9	natural	language Y(c=1)	processing Y(c=2)	and Xκ	machine Y(c=3)	learning	is	fun	and	exciting	#9
#10	natural	language Y(c=1)	processing Y(c=2)	and Xκ	machine Y(c=3)	learning	is	fun	and	exciting	#10



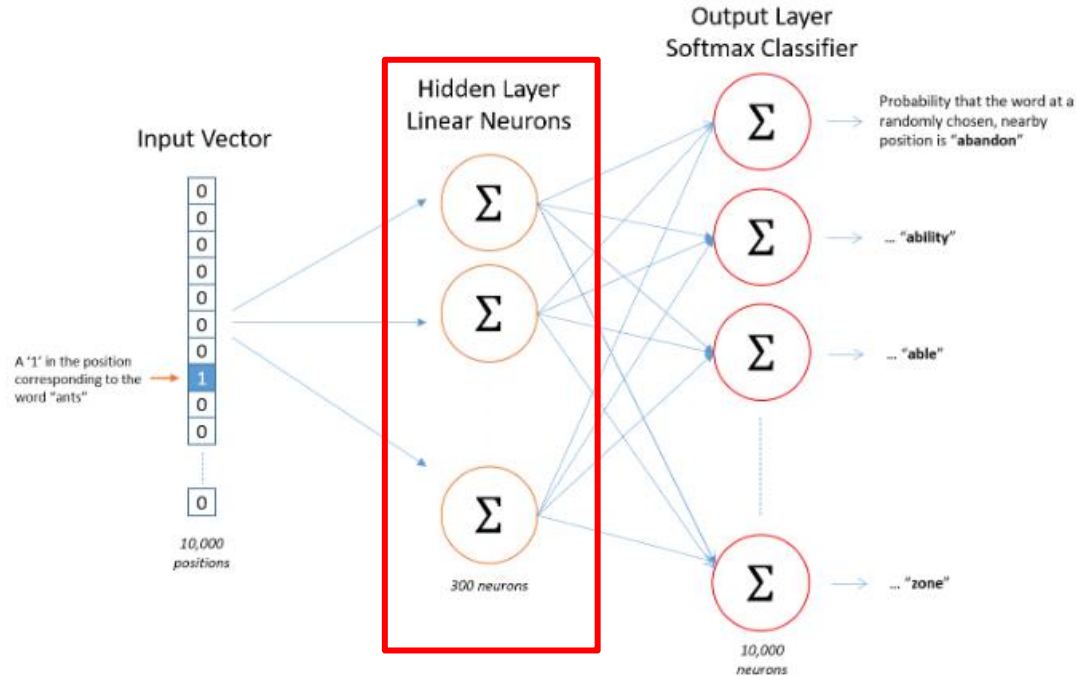
Quick quiz: are language models (large or otherwise), supervised models?

Quick quiz: are language models (large or otherwise), supervised models?

**UNSUPERVISED!**


This allows for big AI companies with lots of compute resources to scale almost indefinitely, because no human labor (with its time, cost, and mistakes) is required!

# How do we get these vectors?



Source: McCormick ML

# What do we do when we pop out the layer?



Let's play with word embeddings!

# **Deep Learning Solutions for NLP**

# Classification Review

- You have a dataset of samples consisting of:
  - $x_i$  inputs, e.g. words (indices to an embedding lookup table or vectors), sentences, documents, etc.
  - $y_i$  outputs which are *labels* of a certain number of classes (binary or multi-class)
- Example:
  - Sentiment analysis
  - Intents in conversational AI
- Architecture
  - Various flavors of deep learning (we'll go over several in this class) topped off by a softmax layer
  - Generally we use cross-entropy loss as our objective function


# Classification Review





# Datasets

- We split our models into, minimally, training and test sets
- You train your model on the training set (normally 70%ish of your data) and then test that same model on your test set
- Preferably train, dev, and test sets -- you train on the training set, then test on dev set iteratively until either you reach your number of epochs or performance ceases to improve, then test is set aside for the “final” numbers on your model
- Ideal situation is thousands of examples with a good amount of variance
- Preferably we want a huge training set and large test set that are “balanced” – that is, we have the same proportion of classes in both the training and the dev and test set



# Softmax

- You can think of it like an activation function (a la the sigmoid or the ReLU) for the final (output) layer of your network
- Normalizes the output of classifier to be some probability for each class, where they all add up to 1:

$$\sigma(x_j) = \frac{e^{x_j}}{\sum_i e^{x_i}}$$

Sample output:

$$\begin{bmatrix} 0.25 \\ 0.21 \\ 0.9 \\ 0.45 \end{bmatrix}$$

- This makes it easier to choose the  $y$  with the maximum probability

# Softmax Example

Original output:

$$\sigma(x_j) = \frac{e^{x_j}}{\sum_i e^{x_i}}$$

$$\begin{bmatrix} 3.0 \\ 5.0 \\ 2.0 \end{bmatrix}$$

$$e^3 + e^5 + e^2$$

$$= 20.0855 + 148.4132 + 7.3891$$

$$= 175.8878$$

# Softmax Example

Original output:

$$\begin{bmatrix} 3.0 \\ 5.0 \\ 2.0 \end{bmatrix}$$

$$20.0855/175.8878 =$$

$$148.4132/175.8878 =$$


$$7.3891/175.8878 =$$

$$\sigma(x_j) = \frac{e^{x_j}}{\sum_i e^{x_i}}$$

$$\begin{bmatrix} 0.12 \\ 0.84 \\ 0.04 \end{bmatrix}$$

# Gradient Descent Review (Hopefully)


- How to find the minimum of your loss function
- Uses the derivative of the loss function with respect to your parameters to take a step in a direction towards the minimum
- How most machine learning algorithms work
- The rate at which you move down this slope is called the learning rate



Picture from Data Science Central, <https://www.datasciencecentral.com/profiles/blogs/alternatives-to-the-gradient-descent-algorithm>

# Cross Entropy Loss

- Cross-Entropy Loss is the loss function we're trying to minimize in classification
- We're attempting to maximize the probability of the correct class  $y$
- Cross entropy loss measures the deviation of the predicted probability from the actual probability



Cross-Entropy Loss

$$-\sum_{c=1}^M y_{o,c} \log(p_{o,c})$$

## Cross-Entropy Loss

$$-\sum_{c=1}^M y_{o,c} \log(p_{o,c})$$

Output of softmax:

Correct label:

$$\begin{bmatrix} 0.12 \\ 0.84 \\ 0.04 \end{bmatrix}$$

$$\begin{bmatrix} 0 \\ 1 \\ 0 \end{bmatrix}$$

$$\begin{aligned} & -(1 \times \log(0.84)) \\ & = 0.174353387 \end{aligned}$$



## Cross-Entropy Loss

$$-\sum_{c=1}^M y_{o,c} \log(p_{o,c})$$

Output of softmax:

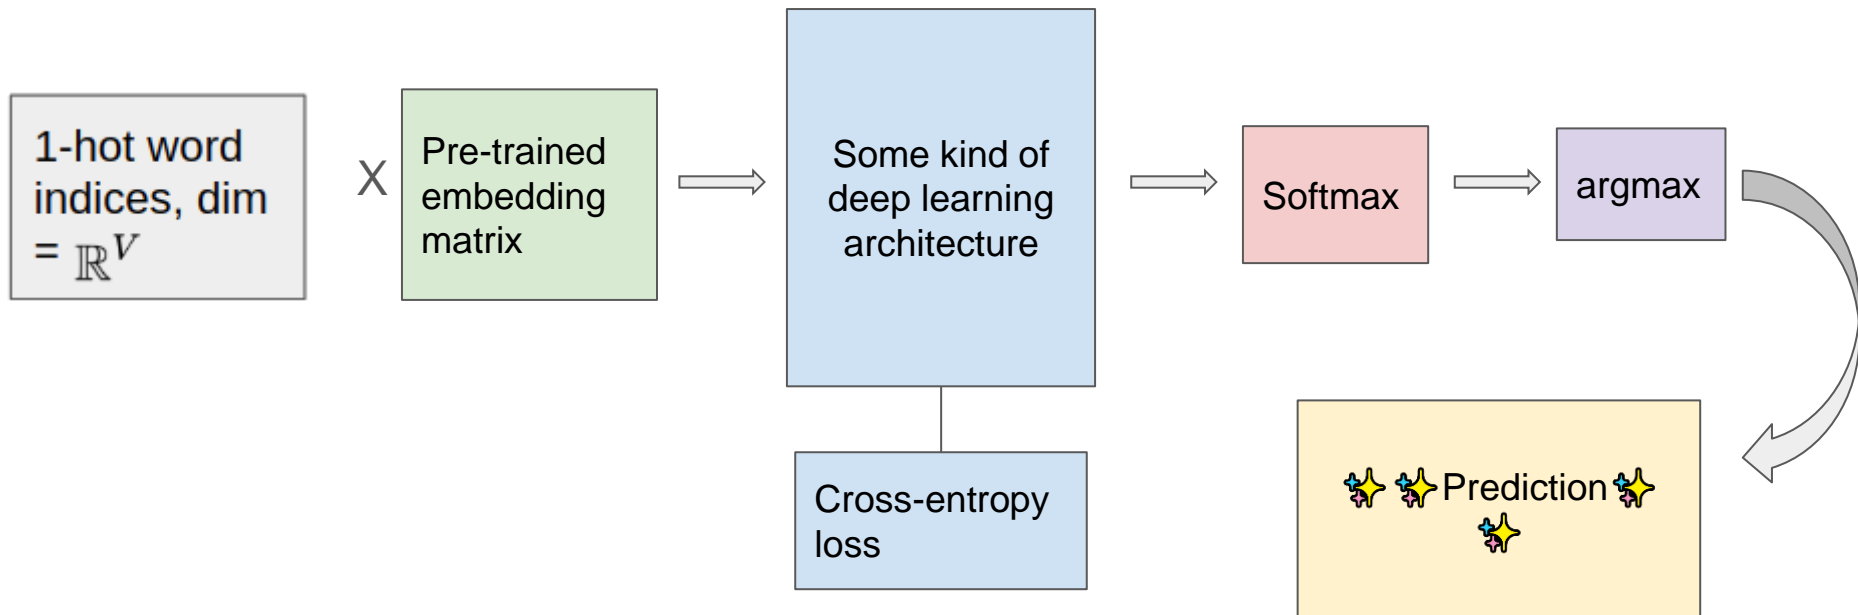
Correct label:

$$\begin{bmatrix} 0.84 \\ 0.12 \\ 0.04 \end{bmatrix}$$


$$\begin{bmatrix} 0 \\ 1 \\ 0 \end{bmatrix}$$

$$\begin{aligned} & -(1 \times \log(0.12)) \\ &= 2.12026353 \end{aligned}$$

# Putting it together



# Putting it together



# Practical Tips for Classification in the Workplace

- Labelers for supervised tasks are an incredibly precious resource
- Many companies outsource to countries like the Philippines to make the labeling less expensive
- If you have the good fortune of working with internal labelers, domain knowledge is more important than labeling experience (in my experience)
- **Make their lives easier whenever you can with unsupervised methods**

# How to Make Your Labelers' Lives Easier (Yours Too)

- Cluster the data to help them make sense of it and break it into chunks (very helpful for higher-dimensional classifiers like intent classifiers)
- Determine what else you can use to filter - e.g. sentiment for something like hate speech or bias detection
- Use simple anomaly detection methods to find likely errors to send back to them
- Use even simpler similarity metrics to find less-frequent examples in a huge haystack
- Be creative! Keep the data representative and pay attention to lower-resource samples


Let's do some clustering!

# Deep Learning Architectures for NLP

- (Vanilla) Recurrent Neural Networks (RNNs)
- Long Short-Term Memory (LSTMs)
- CNNs - they're not just for image processing anymore!
- Transformers
  - BERT, GPT

# Language Modeling with Recurrent Neural Networks (RNNs)

- Words and characters are *\*not\** independent in language, therefore they should not be treated as independent in your neural net
- We'll be revisiting language modeling






# Language Modeling with Recurrent Neural Networks (RNNs)

- Previous methods (n-gram language modeling, see further reading) created incoherent sentences
- Consider the following sentence:

“today the price of gold per ton , while production of shoe lasts and shoe industry , the bank intervened just after it considered and rejected an imf demand to rebuild depleted european stocks , sept 30 end primary 76 cts a share .”

# Language Modeling with Recurrent Neural Networks (RNNs)

- RNNs allow us to use a “memory” for the information that came before



$$s_t = f(Ux_t + Ws_{t-1})$$

# Language Modeling with Recurrent Neural Networks (RNNs)

- To train
  - Get large corpus of text data with a sequence of words  $x_1, \dots, x_T$
  - Predict probability of distribution for each word, given words so far (classification with large number of classes)
  - Loss function is the cross entropy between the predicted next word and the true next word (one-hot vector of class index)
  - Average all of these values to get the overall loss for the training set

# Language Modeling with Recurrent Neural Networks (RNNs)

- RNNs can process any length of input, the model size doesn't increase for longer input
- However, RNNs are slow to compute -- they are difficult to parallelize because the inputs are dependent on one another
- Although they hold some memory, in practice it is difficult to access information from many steps back (5-10 steps)
- Vanishing gradient problem (more on this in a second)

# Language Modeling with Recurrent Neural Networks (RNNs)

PANDARUS:

Alas, I think he shall be come approached and the day

When little strain would be attain'd into being never fed,

And who is but a chain and subjects of his death,

I should not sleep

# What are RNNs capable of?

Clardic Fug	112	113	84
Snowbonk	201	199	165
Catbabel	97	93	68
Bunflow	190	174	155
Ronching Blue	121	114	125
Bank Butt	221	196	199
Caring Tan	171	166	170
Stargoon	233	191	141
Sink	176	138	110
Stummy Beige	216	200	185
Dorkwood	61	63	66
Flower	178	184	196

Sand Dan	201	172	143
Grade Bat	48	94	83
Light Of Blast	175	150	147
Grass Bat	176	99	108
Sindis Poop	204	205	194
Dope	219	209	179
Testing	156	101	106
Stoner Blue	152	165	159
Burple Simp	226	181	132
Stanky Bean	197	162	171
Turdly	190	164	116

Source: [AI Weirdness](#)

# Drawbacks of RNNs

Consider the following word prediction task:


“When she tried to print her tickets, she found that the printer was out of toner. She went to the stationery store to buy more toner. It was very overpriced. After installing the toner into the printer, she finally printed her \_\_\_\_\_”

# Drawbacks of RNNs

- We get something called the Vanishing Gradient Problem
- The gradient can be viewed as a measure of the effect of the past on the future
- If the gradient becomes too small, it's impossible to tell whether there's no dependency between two distant time steps or we have the wrong parameters to capture the dependency between two distant time steps
- This is an artifact of using the chain rule for backpropagation

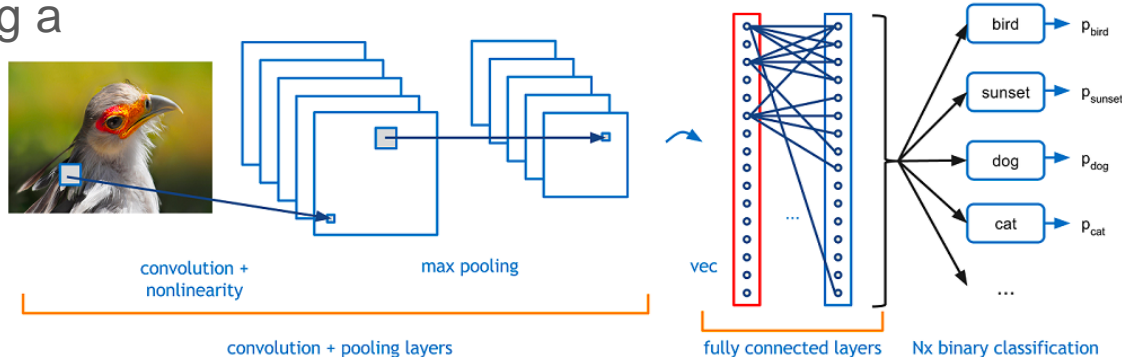


# Drawbacks of RNNs




# CNN Review


- 2-dim image with sliding window
- At each point you create a convolved value by applying a kernel or filter to the output
- You apply some kind of pooling to the outputs to squash the dimensions
- You then use the smaller-dimensional layers to feed as input to your fully-connected layers



# CNNs for NLP



# CNNs for NLP



Now for one more thing...

What things did you pay attention to when you were trying to be a language model? What was intuitive? What was challenging?